

Week5_Assignment

February 20, 2024

```
[1]: import pandas as pd
from pycaret.classification import setup, compare_models, predict_model, \
    save_model, load_model
import pickle
from IPython.display import Code
```

0.1 Load data

```
[3]: df = pd.read_csv("prepped_churn_data.csv")
df
```

```
[3]:
```

	tenure	MonthlyCharges	TotalCharges	Churn	MonthlyCharges_log	\
0	1	29.85	29.85	0	3.396185	
1	34	56.95	1889.50	0	4.042174	
2	2	53.85	108.15	1	3.986202	
3	45	42.30	1840.75	0	3.744787	
4	2	70.70	151.65	1	4.258446	
...	
7027	24	84.80	1990.50	0	4.440296	
7028	72	103.20	7362.90	0	4.636669	
7029	11	29.60	346.45	0	3.387774	
7030	4	74.40	306.60	1	4.309456	
7031	66	105.65	6844.50	0	4.660132	

	TotalCharges_Tenure_Ratio	MonthlyCharges_to_TotalCharges_Ratio	\
0	29.850000	1.000000	
1	55.573529	0.030140	
2	54.075000	0.497920	
3	40.905556	0.022980	
4	75.825000	0.466205	
...	
7027	82.937500	0.042602	
7028	102.262500	0.014016	
7029	31.495455	0.085438	
7030	76.650000	0.242661	
7031	103.704545	0.015436	

	Bank transfer (automatic)	Credit card (automatic)	Electronic check	\
--	---------------------------	-------------------------	------------------	---

0		0		0		0
1		0		0		1
2		0		0		1
3		1		0		1
4		0		0		0
...	
7027		0		0		1
7028		0		1		1
7029		0		0		0
7030		0		0		1
7031		1		0		1

	Mailed check	Month-to-month	One year	Two year
0	0	0	0	0
1	1	1	1	0
2	1	0	0	0
3	0	1	1	0
4	0	0	0	0
...
7027	1	1	1	0
7028	0	1	1	0
7029	0	0	0	0
7030	1	0	0	0
7031	0	1	0	1

[7032 rows x 14 columns]

0.2 initialize auto ML environment

```
[4]: automl_setup = setup(df, target='Churn')
```

```
<pandas.io.formats.style.Styler at 0x7f263f095710>
```

The output summarizes the setup information for the PyCaret auto ML environment.

Session id: 7041 - A unique identifier for the PyCaret session.

Target: Churn - The target variable for the classification task is Churn.

Target type: Binary - The target variable is binary, indicating a binary classification task (Churn or no Churn).

Original data shape: (7032, 14) - The original dataset has 7032 rows and 14 columns.

Transformed data shape: (7032, 14) - The transformed dataset after preprocessing remains the same size as the original dataset.

Transformed train set shape: (4922, 14) - The training set after preprocessing contains 4922 samples.

Transformed test set shape: (2110, 14) - The test set after preprocessing contains 2110 samples.

Numeric features: 13 - There are 13 numeric features in the dataset.

Preprocess: True - The data has been preprocessed.

Imputation type: simple - Simple imputation method has been used for handling missing values.

Numeric imputation: mean - Mean imputation has been applied to numeric features.

Categorical imputation: mode - Mode imputation has been applied to categorical features.

Fold Generator: StratifiedKFold - Stratified K-Fold cross-validation is used during model training.

Number: 10 - 10 folds are used in cross-validation.

CPU Jobs: -1 - The number of CPU jobs is set to -1, allowing PyCaret to utilize all available CPUs.

Use GPU: False - GPU acceleration is not utilized for model training.

Log Experiment: False - Logging of the experiment is turned off.

Experiment Name: clf-default-name - The default name for the classification experiment is 'clf-default-name'.

USI: cc2a - A unique identifier for the experiment setup.

```
[5]: automl_type = type(automl_setup)
      automl_type
```

```
[5]: pycaret.classification.oop.ClassificationExperiment
```

0.3 Compare and select best model

```
[6]: best_model = compare_models()
```

```
<IPython.core.display.HTML object>
```

```
<pandas.io.formats.style.Styler at 0x7f263ea74050>
```

```
<IPython.core.display.HTML object>
```

This output summarizes the performance metrics of various machine learning models trained on the prepped churn dataset, including accuracy, area under the curve (AUC), recall, precision, F1 score, Kappa, Matthews correlation coefficient (MCC), and training time in seconds.

The best performing model based on accuracy:

The Ridge Classifier achieved the highest accuracy of 79.42% followed closely by Logistic Regression accuracy of 79.38%. LDA is another high performer with an accuracy of 79.36%.

Interpreting the results

Accuracy: Indicates the proportion of correctly classified instances out of the total instances.

AUC: Represents the area under the receiver operating characteristic (ROC) curve, which measures the model's ability to distinguish between classes.

Recall: Denotes the proportion of actual positive cases that were correctly identified by the model.

Precision: Indicates the proportion of positive identifications that were actually correct.

F1 Score: Harmonic mean of precision and recall, providing a balance between the two metrics.

Kappa: Measures the agreement between predicted and actual classifications, considering the possibility of the agreement occurring by chance.

MCC (Matthews Correlation Coefficient): Another measure of the quality of binary classifications, considering both false positives and false negatives.

Training Time (TT): Indicates the time taken by each model to train on the dataset.

```
[7]: best_model_info = best_model
      best_model_info
```

```
[7]: RidgeClassifier(alpha=1.0, class_weight=None, copy_X=True, fit_intercept=True,
                    max_iter=None, positive=False, random_state=7041, solver='auto',
                    tol=0.0001)
```

0.4 Select specific rows

```
[8]: selected_rows = df.iloc[400:415]
      selected_rows
```

```
[8]:
```

	tenure	MonthlyCharges	TotalCharges	Churn	MonthlyCharges_log	\
400	32	19.75	624.15	0	2.983153	
401	11	20.05	237.70	0	2.998229	
402	69	99.45	7007.60	1	4.599655	
403	68	55.90	3848.80	0	4.023564	
404	20	19.70	419.40	0	2.980619	
405	72	19.80	1468.75	0	2.985682	
406	60	95.40	5812.00	0	4.558079	
407	32	93.95	2861.45	0	4.542763	
408	1	19.90	19.90	1	2.990720	
409	1	19.60	19.60	1	2.975530	
410	3	81.35	233.70	1	4.398761	
411	46	24.45	1066.15	0	3.196630	
412	29	74.95	2149.05	0	4.316821	
413	51	87.35	4473.00	0	4.469923	
414	48	70.65	3545.05	0	4.257738	

	TotalCharges_Tenure_Ratio	MonthlyCharges_to_TotalCharges_Ratio	\
400	19.504687	0.031643	
401	21.609091	0.084350	
402	101.559420	0.014192	
403	56.600000	0.014524	
404	20.970000	0.046972	
405	20.399306	0.013481	
406	96.866667	0.016414	
407	89.420312	0.032833	
408	19.900000	1.000000	
409	19.600000	1.000000	

410	77.900000	0.348096
411	23.177174	0.022933
412	74.105172	0.034876
413	87.705882	0.019528
414	73.855208	0.019929

	Bank transfer (automatic)	Credit card (automatic)	Electronic check \
400	1	0	1
401	0	1	1
402	0	1	1
403	1	0	1
404	0	0	1
405	0	1	1
406	1	0	1
407	0	0	1
408	0	0	1
409	0	0	1
410	0	0	0
411	1	0	1
412	0	0	0
413	0	0	0
414	1	0	1

	Mailed check	Month-to-month	One year	Two year
400	0	1	1	0
401	0	1	1	0
402	0	0	0	0
403	0	1	1	0
404	1	1	0	1
405	0	1	0	1
406	0	1	1	0
407	1	0	0	0
408	1	0	0	0
409	1	0	0	0
410	0	0	0	0
411	0	1	1	0
412	0	0	0	0
413	0	0	0	0
414	0	0	0	0

0.5 Utilize best model to predict churn

```
[9]: predict_model(best_model, selected_rows)
```

```
<pandas.io.formats.style.Styler at 0x7f263e72c250>
```

```
[9]:      tenure  MonthlyCharges  TotalCharges  MonthlyCharges_log  \
400      32      19.750000      624.150024      2.983154
401      11      20.049999      237.699997      2.998229
402      69      99.449997      7007.600098      4.599655
403      68      55.900002      3848.800049      4.023564
404      20      19.700001      419.399994      2.980619
405      72      19.799999      1468.750000      2.985682
406      60      95.400002      5812.000000      4.558079
407      32      93.949997      2861.449951      4.542763
408       1      19.900000      19.900000      2.990720
409       1      19.600000      19.600000      2.975530
410       3      81.349998      233.699997      4.398761
411      46      24.450001      1066.150024      3.196630
412      29      74.949997      2149.050049      4.316821
413      51      87.349998      4473.000000      4.469923
414      48      70.650002      3545.050049      4.257738
```

```
      TotalCharges_Tenure_Ratio  MonthlyCharges_to_TotalCharges_Ratio  \
400      19.504688      0.031643
401      21.609091      0.084350
402     101.559418      0.014192
403      56.599998      0.014524
404      20.969999      0.046972
405      20.399305      0.013481
406      96.866669      0.016414
407      89.420311      0.032833
408      19.900000      1.000000
409      19.600000      1.000000
410      77.900002      0.348096
411      23.177174      0.022933
412      74.105171      0.034876
413      87.705879      0.019528
414      73.855209      0.019929
```

```
      Bank transfer (automatic)  Credit card (automatic)  Electronic check  \
400      1      0      1
401      0      1      1
402      0      1      1
403      1      0      1
404      0      0      1
405      0      1      1
406      1      0      1
407      0      0      1
408      0      0      1
409      0      0      1
410      0      0      0
411      1      0      1
```

412	0	0	0
413	0	0	0
414	1	0	1

	Mailed check	Month-to-month	One year	Two year	Churn	prediction_label
400	0	1	1	0	0	0
401	0	1	1	0	0	0
402	0	0	0	0	1	0
403	0	1	1	0	0	0
404	1	1	0	1	0	0
405	0	1	0	1	0	0
406	0	1	1	0	0	0
407	1	0	0	0	0	0
408	1	0	0	0	1	0
409	1	0	0	0	1	0
410	0	0	0	0	1	1
411	0	1	1	0	0	0
412	0	0	0	0	0	0
413	0	0	0	0	0	0
414	0	0	0	0	0	0

Metrics

Model: Ridge Classifier Accuracy: 80% AUC: 62.5% Recall: 25% Precision: 100% F1 Score: 40%
Kappa: 32.84% MCC: 44.32%

While the model exhibits high precision, suggesting it correctly identifies churn when it occurs, its recall is quite low, indicating it misses many actual churn instances.

0.6 Incorrect predictions

```
[10]: predicted_rows = predict_model(best_model, selected_rows)
      incorrect_predictions = (predicted_rows['Churn'] !=
      ↪ predicted_rows['prediction_label']).sum()

      print("Incorrect Predictions:", incorrect_predictions)
```

<pandas.io.formats.style.Styler at 0x7f263fec20d0>

Incorrect Predictions: 3

Out of the total predictions made, the model was incorrect in predicting the churn status of 3 customers.

0.7 Save model

```
[11]: save_model(best_model, 'ridge')
```

Transformation Pipeline and Model Successfully Saved

```
[11]: (Pipeline(memory=Memory(location=None),
              steps=[('numerical_imputer',
                      TransformerWrapper(exclude=None,
                                         include=['tenure', 'MonthlyCharges',
                                                'TotalCharges',
                                                'MonthlyCharges_log',
                                                'TotalCharges_Tenure_Ratio',
                                                'MonthlyCharges_to_TotalCharges_Ratio',
                                                'Bank transfer (automatic)',
                                                'Credit card (automatic)',
                                                'Electronic check', 'Mailed
check',
                                                'Month-to-month', 'One year',
                                                'Two y...

strategy='most_frequent',
verbose='deprecated'))),
      ('clean_column_names',
       TransformerWrapper(exclude=None, include=None,
                           transformer=CleanColumnNames(match='[\\]\\\\[\\,\\\\{\\\\}\\\\"\\\\:]+'))),
      ('trained_model',
       RidgeClassifier(alpha=1.0, class_weight=None, copy_X=True,
                       fit_intercept=True, max_iter=None,
                       positive=False, random_state=7041,
                       solver='auto', tol=0.0001))),
      verbose=False),
      'ridge.pkl')
```

0.8 Use pickle to save and load the model (serialization and deserialization)

```
[14]: with open('ridge_model.pk', 'wb') as f:
      pickle.dump(best_model, f)
```

```
[16]: with open('ridge_model.pk', 'rb') as f:
      loaded_model = pickle.load(f)
```

0.9 Create new data

```
[17]: new_data = selected_rows.copy()
      new_data.drop('Churn', axis=1, inplace=True)
      new_data.to_csv('new_churn_data.csv', index=False)
```

0.10 Make predictions for churn on the loaded data

```
[18]: loaded_model.predict(new_data)
```

```
[18]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0], dtype=int8)
```



```
[19]: loaded_ridge = load_model('ridge')
      predict_model(loaded_ridge, new_data)
```

Transformation Pipeline and Model Successfully Loaded

<IPython.core.display.HTML object>

```
[19]:      tenure  MonthlyCharges  TotalCharges  MonthlyCharges_log  \
400      32      19.750000      624.150024      2.983154
401      11      20.049999      237.699997      2.998229
402      69      99.449997      7007.600098      4.599655
403      68      55.900002      3848.800049      4.023564
404      20      19.700001      419.399994      2.980619
405      72      19.799999      1468.750000      2.985682
406      60      95.400002      5812.000000      4.558079
407      32      93.949997      2861.449951      4.542763
408       1      19.900000      19.900000      2.990720
409       1      19.600000      19.600000      2.975530
410       3      81.349998      233.699997      4.398761
411      46      24.450001      1066.150024      3.196630
412      29      74.949997      2149.050049      4.316821
413      51      87.349998      4473.000000      4.469923
414      48      70.650002      3545.050049      4.257738
```

```
      TotalCharges_Tenure_Ratio  MonthlyCharges_to_TotalCharges_Ratio  \
400      19.504688      0.031643
401      21.609091      0.084350
402     101.559418      0.014192
403      56.599998      0.014524
404      20.969999      0.046972
405      20.399305      0.013481
406      96.866669      0.016414
407      89.420311      0.032833
408      19.900000      1.000000
409      19.600000      1.000000
410      77.900002      0.348096
411      23.177174      0.022933
412      74.105171      0.034876
413      87.705879      0.019528
414      73.855209      0.019929
```

```
      Bank transfer (automatic)  Credit card (automatic)  Electronic check  \
400      1      0      1
401      0      1      1
402      0      1      1
403      1      0      1
404      0      0      1
405      0      1      1
```

406	1	0	1
407	0	0	1
408	0	0	1
409	0	0	1
410	0	0	0
411	1	0	1
412	0	0	0
413	0	0	0
414	1	0	1

	Mailed check	Month-to-month	One year	Two year	prediction_label
400	0	1	1	0	0
401	0	1	1	0	0
402	0	0	0	0	0
403	0	1	1	0	0
404	1	1	0	1	0
405	0	1	0	1	0
406	0	1	1	0	0
407	1	0	0	0	0
408	1	0	0	0	0
409	1	0	0	0	0
410	0	0	0	0	1
411	0	1	1	0	0
412	0	0	0	0	0
413	0	0	0	0	0
414	0	0	0	0	0

0.11 Python Module to predict churn

```
[20]: Code('predict_churn.py')
```

```
[20]: import pandas as pd
from pycaret.classification import predict_model, load_model

def predict_churn():
    df = pd.read_csv('new_churn_data.csv')
    model = load_model('ridge')
    predictions = predict_model(model, df)
    predictions.rename({'prediction_label': 'Churn_prediction'}, axis=1,
inplace=True)
    predictions['Churn_prediction'].replace({1: 'Churn', 0: 'No Churn'},
inplace=True)
    return predictions['Churn_prediction']

# Call the function and print the predictions
print(predict_churn())
```

```
[21]: %run predict_churn.py
```

Transformation Pipeline and Model Successfully Loaded

<IPython.core.display.HTML object>

```
0      No Churn
1      No Churn
2      No Churn
3      No Churn
4      No Churn
5      No Churn
6      No Churn
7      No Churn
8      No Churn
9      No Churn
10     Churn
11     No Churn
12     No Churn
13     No Churn
14     No Churn
```

Name: Churn_prediction, dtype: object

The output indicates the churn predictions for each customer in the dataset. Each entry in the output corresponds to a customer, and it shows whether the model predicts that the customer will churn or not churn.

Necessary libraries and functions were imported for this process, which involves building a churn prediction model using PyCaret, a Python library for automating machine learning workflows.

We successfully achieved the following,

Loaded and prepared the churn data.

Set up an auto ML environment and compared classification models.

Selected the best-performing model which was Ridge Classifier

Predicted the churn status for 15 specific rows of data using the selected model.

Saved the best-performing model to a file using PyCaret's `save_model` function.

Serialized and deserialized the model using pickle.

Predicted the churn status for new data using both the loaded model and PyCaret's `load_model` function